

```
In [2]: #importing pandas library
import pandas as pd
#reading the dataset
data = pd.read_csv( "eps_AA.csv")
data.head()
```

```
Out[2]:
```

	TICKER	CNAME	ACTDATS	ESTIMATOR	ANALYS	FPI	MEASURE	VALUE	FPEDATS	REVDATS
0	AA	ALCOA	19950112	118.0	288.0	1.0	EPS	2.2313	19951231.0	19950228.0
1	AA	ALCOA INC.	20111005	2488.0	18082.0	1.0	EPS	2.7600	20111231.0	20111011.0
2	AA	ALCOA INC.	20111006	1267.0	73367.0	1.0	EPS	2.9400	20111231.0	20111006.0
3	AA	ALCOA INC.	20111010	11.0	107166.0	1.0	EPS	3.5700	20111231.0	20111010.0
4	AA	ALCOA INC.	20111011	118.0	112989.0	1.0	EPS	2.6100	20111231.0	20111011.0

```
In [3]: # Display the first few rows of the dataset
# to understand its structure
data.head(), data.info(), data.describe(), data.shape
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2845 entries, 0 to 2844
Data columns (total 16 columns):
#   Column          Non-Null Count  Dtype
---  -
0   TICKER          2845 non-null   object
1   CNAME           2845 non-null   object
2   ACTDATS         2845 non-null   int64
3   ESTIMATOR       2845 non-null   float64
4   ANALYS          2845 non-null   float64
5   FPI             2845 non-null   float64
6   MEASURE         2845 non-null   object
7   VALUE          2845 non-null   float64
8   FPEDATS         2845 non-null   float64
9   REVDATS         2845 non-null   float64
10  REVTIMS         2845 non-null   object
11  ANNDATS         2845 non-null   float64
12  ANNTIMS         2845 non-null   object
13  ACTUAL          2801 non-null   float64
14  ANNDATS_ACT     2801 non-null   float64
15  ANNTIMS_ACT     2801 non-null   object
dtypes: float64(9), int64(1), object(6)
memory usage: 355.8+ KB
```

```

Out[3]: (  TICKER      CNAME    ACTDATS  ESTIMATOR    ANALYS  FPI  MEASURE    VALUE  \
0      AA      ALCOA    19950112      118.0      288.0  1.0      EPS  2.2313
1      AA  ALCOA INC.    20111005      2488.0    18082.0  1.0      EPS  2.7600
2      AA  ALCOA INC.    20111006      1267.0    73367.0  1.0      EPS  2.9400
3      AA  ALCOA INC.    20111010       11.0   107166.0  1.0      EPS  3.5700
4      AA  ALCOA INC.    20111011      118.0   112989.0  1.0      EPS  2.6100

      FPEDATS      REVDATS    REVTIMS      ANNDATS    ANNTIMS  ACTUAL  \
0  19951231.0  19950228.0  13:49:55  19950112.0  11:06:14  3.375
1  20111231.0  20111011.0  10:30:27  20111004.0  7:04:00  2.160
2  20111231.0  20111006.0  19:18:18  20111006.0  17:55:00  2.160
3  20111231.0  20111010.0  16:06:23  20111010.0  15:11:00  2.160
4  20111231.0  20111011.0  10:28:04  20111010.0  20:30:00  2.160

      ANNDATS_ACT  ANNTIMS_ACT
0  19960108.0      18:42:00
1  20120109.0      16:03:00
2  20120109.0      16:03:00
3  20120109.0      16:03:00
4  20120109.0      16:03:00 ,
None,

      ACTDATS      ESTIMATOR      ANALYS      FPI      VALUE  \
count  2.845000e+03  2845.000000  2845.000000  2845.0  2845.000000
mean    2.007413e+07   755.387346  62086.166960    1.0    3.520233
std     6.990649e+04  1009.991771  49944.681507    0.0    2.764763
min     1.995011e+07   11.000000   127.000000    1.0   -6.330000
25%     2.002011e+07   109.000000  18096.000000    1.0    1.500000
50%     2.007113e+07   192.000000  73367.000000    1.0    3.375000
75%     2.013071e+07  1110.000000  107166.000000    1.0    5.100000
max     2.021121e+07  4280.000000  194868.000000    1.0   11.790000

      FPEDATS      REVDATS      ANNDATS      ACTUAL  ANNDATS_ACT
count  2.845000e+03  2.845000e+03  2.845000e+03  2801.000000  2.801000e+03
mean    2.007439e+07  2.007533e+07  2.007409e+07    3.115243  2.008112e+07
std     6.973694e+04  6.892611e+04  6.995179e+04    2.399600  6.813478e+04
min     1.995123e+07  1.995022e+07  1.995011e+07   -2.400000  1.996011e+07
25%     2.002123e+07  2.002023e+07  2.002011e+07    1.620000  2.002011e+07
50%     2.007123e+07  2.007122e+07  2.007113e+07    2.760000  2.008011e+07
75%     2.013123e+07  2.013072e+07  2.013071e+07    4.380000  2.014011e+07
max     2.021123e+07  2.021122e+07  2.021121e+07    8.880000  2.021020e+07 ,
(2845, 16))

```

## Data Cleaning

```

In [4]: #Checking null values in dataset
null_counts = data.isnull().sum()
print(null_counts)

```

```

TICKER      0
CNAME       0
ACTDATS     0
ESTIMATOR   0
ANALYS      0
FPI         0
MEASURE     0
VALUE       0
FPEDATS     0
REVDATS     0
REVTIMS     0
ANNDATS     0
ANNTIMS     0
ACTUAL      44
ANNDATS_ACT 44
ANNTIMS_ACT 44
dtype: int64

```

```

In [5]: # Remove rows with missing values
clean_data = data.dropna()

```

```

In [6]: null_values = clean_data.isnull().sum()
# Display the results
print("Null Values in Dataset:\n", null_values)

```

Null Values in Dataset:

```

TICKER      0
CNAME       0
ACTDATS     0
ESTIMATOR   0
ANALYS      0
FPI         0
MEASURE     0
VALUE       0
FPEDATS     0
REVDATS     0
REVTIMS     0
ANNDATS     0
ANNTIMS     0
ACTUAL      0
ANNDATS_ACT 0
ANNTIMS_ACT 0
dtype: int64

```

```

In [7]: # Convert the 'REVDATS' column to datetime format for accurate sorting
clean_data['REVDATS'] = pd.to_datetime(clean_data['REVDATS'], format='%Y%m%d')
clean_data['FPEDATS'] = pd.to_datetime(clean_data['FPEDATS'], format='%Y%m%d')
clean_data['ANNDATS_ACT'] = pd.to_datetime(clean_data['ANNDATS_ACT'], format='%Y%m%d')

# Sort data by 'ANALYS', 'FPEDATS' (fiscal year), and 'REVDATS' (revision date), and t
clean_data = clean_data.sort_values(by=['ANALYS', 'FPEDATS', 'REVDATS'], ascending=[Tr
clean_data = clean_data.drop_duplicates(subset=['ANALYS', 'FPEDATS'], keep='first')

```

```

C:\Users\YOG\AppData\Local\Temp\ipykernel_22772\3625313825.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    clean_data['REVDATS'] = pd.to_datetime(clean_data['REVDATS'], format='%Y%m%d')
C:\Users\YOG\AppData\Local\Temp\ipykernel_22772\3625313825.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    clean_data['FPEDATS'] = pd.to_datetime(clean_data['FPEDATS'], format='%Y%m%d')
C:\Users\YOG\AppData\Local\Temp\ipykernel_22772\3625313825.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    clean_data['ANNDATS_ACT'] = pd.to_datetime(clean_data['ANNDATS_ACT'], format='%Y%m%d')

```

```

In [8]: # Validate that for each analyst and fiscal year, the date kept is the maximum date
        grouped = clean_data.groupby(['ANALYS', 'FPEDATS'])
        max_dates = grouped['REVDATS'].transform('max')
        assert all(clean_data['REVDATS'] == max_dates), "Not all entries are the most recent p

        # Display the cleaned data
        print(clean_data.head())
        print(clean_data.info())

```

	TICKER	CNAME	ACTDATS	ESTIMATOR	ANALYS	FPI	MEASURE	VALUE	\
1813	AA	ALCOA	19951227	16.0	127.0	1.0	EPS	3.4950	
2071	AA	ALCOA	19961008	16.0	127.0	1.0	EPS	2.7375	
2103	AA	ALCOA	19971224	16.0	127.0	1.0	EPS	3.4050	
1544	AA	ALCOA	19981007	16.0	127.0	1.0	EPS	3.6000	
1432	AA	ALCOA INC.	19991007	16.0	127.0	1.0	EPS	4.0950	

	FPEDATS	REVDATS	REVTIMS	ANNDATS	ANNTIMS	ACTUAL	\
1813	1995-12-31	1998-01-09	17:09:26	19951227.0	10:21:39	3.3750	
2071	1996-12-31	1999-01-08	15:50:40	19961008.0	16:40:06	2.7375	
2103	1997-12-31	2000-01-14	15:31:10	19971224.0	14:01:01	3.3375	
1544	1998-12-31	2000-06-26	16:17:09	19981007.0	13:32:22	3.6525	
1432	1999-12-31	1999-12-02	12:52:09	19991007.0	10:19:11	4.2300	

	ANNDATS_ACT	ANNTIMS_ACT
1813	1996-01-08	18:42:00
2071	1997-01-08	19:04:00
2103	1998-01-08	20:23:00
1544	1999-01-08	8:59:00
1432	2000-01-10	9:09:00

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 539 entries, 1813 to 913
```

```
Data columns (total 16 columns):
```

#	Column	Non-Null Count	Dtype
0	TICKER	539 non-null	object
1	CNAME	539 non-null	object
2	ACTDATS	539 non-null	int64
3	ESTIMATOR	539 non-null	float64
4	ANALYS	539 non-null	float64
5	FPI	539 non-null	float64
6	MEASURE	539 non-null	object
7	VALUE	539 non-null	float64
8	FPEDATS	539 non-null	datetime64[ns]
9	REVDATS	539 non-null	datetime64[ns]
10	REVTIMS	539 non-null	object
11	ANNDATS	539 non-null	float64
12	ANNTIMS	539 non-null	object
13	ACTUAL	539 non-null	float64
14	ANNDATS_ACT	539 non-null	datetime64[ns]
15	ANNTIMS_ACT	539 non-null	object

```
dtypes: datetime64[ns](3), float64(6), int64(1), object(6)
```

```
memory usage: 71.6+ KB
```

```
None
```

## Feature Engineering

```
In [9]: #Calculating past_accuracy
# Convert FPEDATS to datetime and extract the year
clean_data['FPEDATS'] = pd.to_datetime(clean_data['FPEDATS'], format='%Y%m%d', errors=

# Sort the data by analyst, company, and fiscal year
clean_data.sort_values(by=['ANALYS', 'CNAME', 'FPEDATS'], inplace=True)

# Calculate past_accuracy as the difference between last year's VALUE and this year's
clean_data['past_accuracy'] = clean_data.groupby(['ANALYS', 'CNAME'])['VALUE'].shift(1

# Remove rows where 'past_accuracy' is NaN
```

```
clean_data.dropna(subset=['past_accuracy'], inplace=True)
```

```
# Display to check the values
```

```
print(clean_data[['ANALYS', 'CNAME', 'FPEDATS', 'VALUE', 'ACTUAL', 'past_accuracy']].t
```

	ANALYS	CNAME	FPEDATS	VALUE	ACTUAL	past_accuracy
2071	127.0	ALCOA	1996	2.7375	2.7375	0.7575
2103	127.0	ALCOA	1997	3.4050	3.3375	-0.6000
1544	127.0	ALCOA	1998	3.6000	3.6525	-0.2475
1486	127.0	ALCOA INC.	2000	6.3750	5.4450	-1.3500
1786	281.0	ALCOA	1996	4.1250	2.7375	0.6375

```
In [10]: # Ensure both ANNDATS_ACT and ANNDATS are in datetime format
```

```
clean_data['ANNDATS_ACT'] = pd.to_datetime(clean_data['ANNDATS_ACT'], format='%Y%m%d',
clean_data['ANNDATS'] = pd.to_datetime(clean_data['ANNDATS'], format='%Y%m%d', errors=
```

```
# Now calculate the forecast horizon
```

```
clean_data['horizon'] = (clean_data['ANNDATS_ACT'] - clean_data['ANNDATS']).dt.days
```

```
# Check the results
```

```
print(clean_data[['ANNDATS_ACT', 'ANNDATS', 'horizon']].head())
```

	ANNDATS_ACT	ANNDATS	horizon
2071	1997-01-08	1996-10-08	92
2103	1998-01-08	1997-12-24	15
1544	1999-01-08	1998-10-07	93
1486	2001-01-08	2000-04-10	273
1786	1997-01-08	1996-01-31	343

```
In [11]: # Calculate brokerage house size
```

```
clean_data['size'] = clean_data.groupby(['ESTIMATOR', 'CNAME', 'FPEDATS'])['ANALYS'].t
```

```
# Display a few rows to check the 'size' values
```

```
print(clean_data[['ESTIMATOR', 'CNAME', 'FPEDATS', 'ANALYS', 'size']].head())
```

	ESTIMATOR	CNAME	FPEDATS	ANALYS	size
2071	16.0	ALCOA	1996	127.0	1
2103	16.0	ALCOA	1997	127.0	1
1544	16.0	ALCOA	1998	127.0	1
1486	16.0	ALCOA INC.	2000	127.0	1
1786	231.0	ALCOA	1996	281.0	1

```
In [12]: # Calculate analyst experience
```

```
clean_data['experience'] = clean_data.groupby(['ANALYS', 'CNAME'])['FPEDATS'].transform
```

```
# Display a few rows to check the 'experience' values
```

```
print(clean_data[['ANALYS', 'CNAME', 'FPEDATS', 'experience']].head())
```

	ANALYS	CNAME	FPEDATS	experience
2071	127.0	ALCOA	1996	1
2103	127.0	ALCOA	1997	2
1544	127.0	ALCOA	1998	3
1486	127.0	ALCOA INC.	2000	1
1786	281.0	ALCOA	1996	1

```
In [13]: # Validate and display the cleaned data with new features
```

```
print(clean_data[['past_accuracy', 'horizon', 'experience', 'size']].head())
```

```
print(clean_data.info())
```

```

      past_accuracy horizon experience size
2071      0.7575      92          1      1
2103     -0.6000      15          2      1
1544     -0.2475      93          3      1
1486     -1.3500     273          1      1
1786      0.6375     343          1      1
<class 'pandas.core.frame.DataFrame'>
Int64Index: 345 entries, 2071 to 913
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   TICKER                 345 non-null    object
1   CNAME                  345 non-null    object
2   ACTDATS                345 non-null    int64
3   ESTIMATOR              345 non-null    float64
4   ANALYS                 345 non-null    float64
5   FPI                    345 non-null    float64
6   MEASURE                 345 non-null    object
7   VALUE                  345 non-null    float64
8   FPEDATS                345 non-null    int64
9   REVDATS                345 non-null    datetime64[ns]
10  REVTIMS                 345 non-null    object
11  ANNDATS                345 non-null    datetime64[ns]
12  ANNTIMS                 345 non-null    object
13  ACTUAL                  345 non-null    float64
14  ANNDATS_ACT             345 non-null    datetime64[ns]
15  ANNTIMS_ACT             345 non-null    object
16  past_accuracy           345 non-null    float64
17  horizon                 345 non-null    int64
18  size                    345 non-null    int64
19  experience               345 non-null    int32
dtypes: datetime64[ns](3), float64(6), int32(1), int64(4), object(6)
memory usage: 55.3+ KB
None

```

## Descriptive Statistics

```

In [14]: # Assuming 'clean_data' is your cleaned and processed DataFrame
descriptive_stats = clean_data[['MEASURE', 'ACTUAL', 'past_accuracy', 'horizon', 'experience', 'size']]

# Print the descriptive statistics
print(descriptive_stats)

```

	ACTUAL	past_accuracy	horizon	experience	size
count	345.000000	345.000000	345.000000	345.000000	345.000000
mean	3.403797	0.115459	93.286957	2.97971	1.034783
std	2.447899	2.336796	99.483617	2.51094	0.183495
min	-2.400000	-4.770000	0.000000	1.00000	1.000000
25%	2.110000	-1.140000	20.000000	1.00000	1.000000
50%	3.300000	-0.130000	78.000000	2.00000	1.000000
75%	4.620000	1.020000	111.000000	4.00000	1.000000
max	8.880000	7.452000	546.000000	13.00000	2.000000


## Exploratory Data Analysis

```
In [15]: #Downloading necessary libraries
```

```
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
In [16]: # Scatter plot for 'past_accuracy' vs. 'ACTUAL'
```


```
plt.figure(figsize=(10, 6))
sns.scatterplot(x='past_accuracy', y='ACTUAL', data=clean_data)
plt.title('Scatter Plot of Past Accuracy vs. Actual EPS')
plt.xlabel('Past Accuracy')
plt.ylabel('Actual EPS')
plt.show()
```



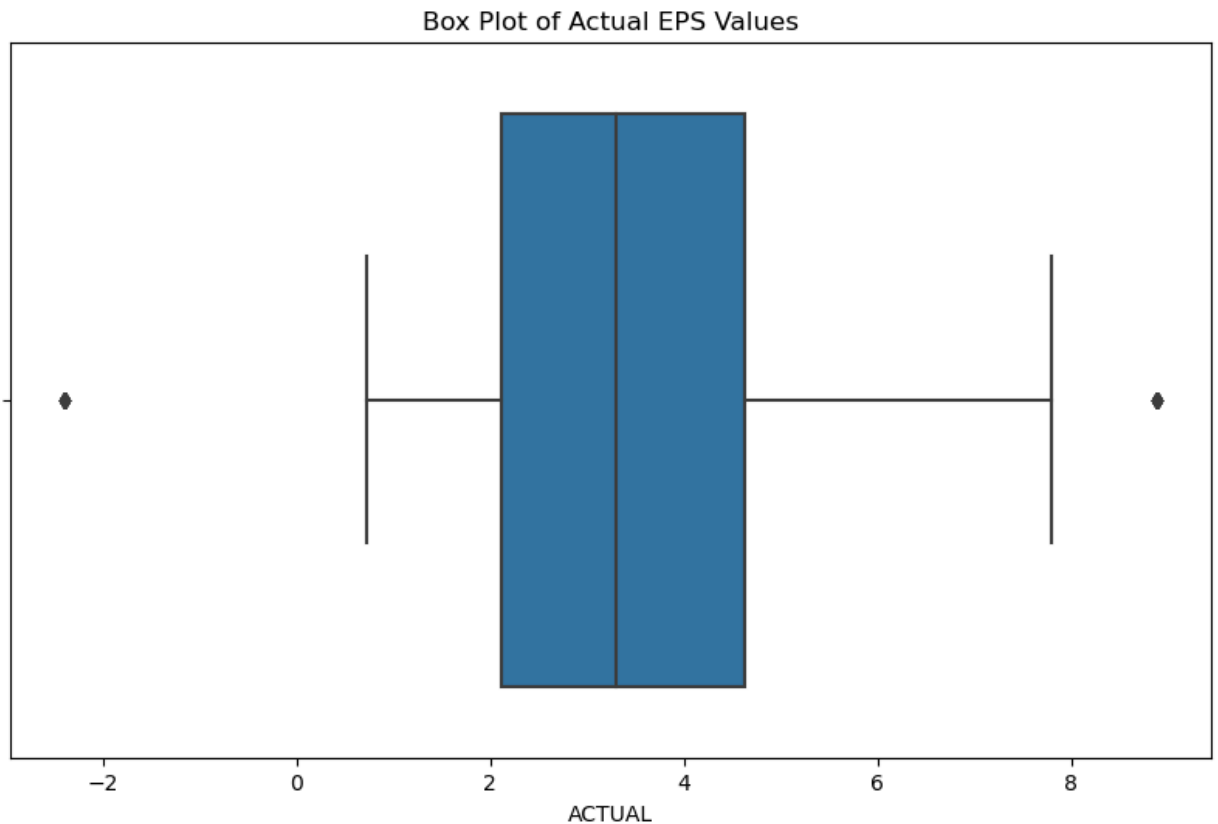
```
In [17]: # Histogram for 'experience'
```

```
plt.figure(figsize=(10, 6))
sns.histplot(clean_data['experience'], kde=True)
plt.title('Distribution of Analyst Experience')
plt.xlabel('Years of Experience')
plt.show()
```



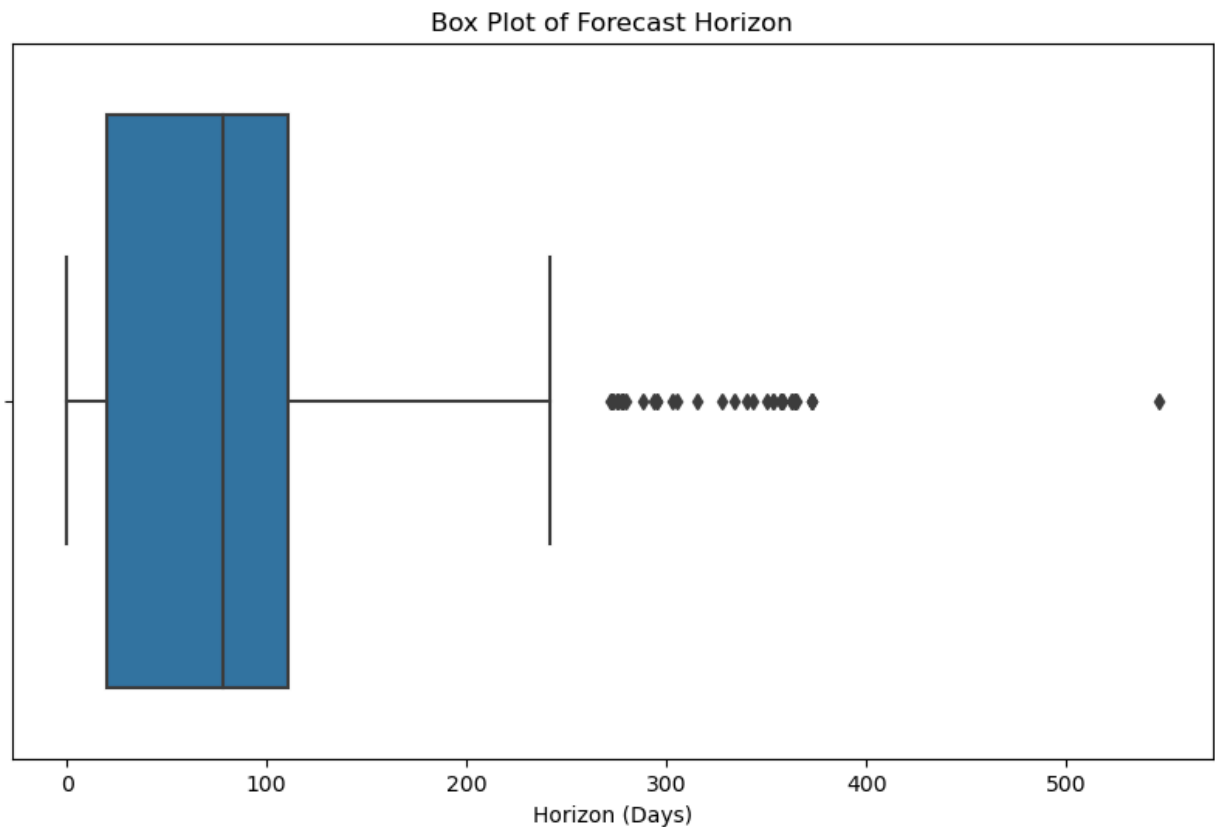


```
In [18]: # Box plot for 'actual'
plt.figure(figsize=(10, 6))
sns.boxplot(x='ACTUAL', data=clean_data)
plt.title('Box Plot of Actual EPS Values')
plt.show()
```



```
In [19]: # box plot for horizon

plt.figure(figsize=(10, 6))
sns.boxplot(x='horizon', data=clean_data)
plt.title('Box Plot of Forecast Horizon')
plt.xlabel('Horizon (Days)')
plt.show()
```

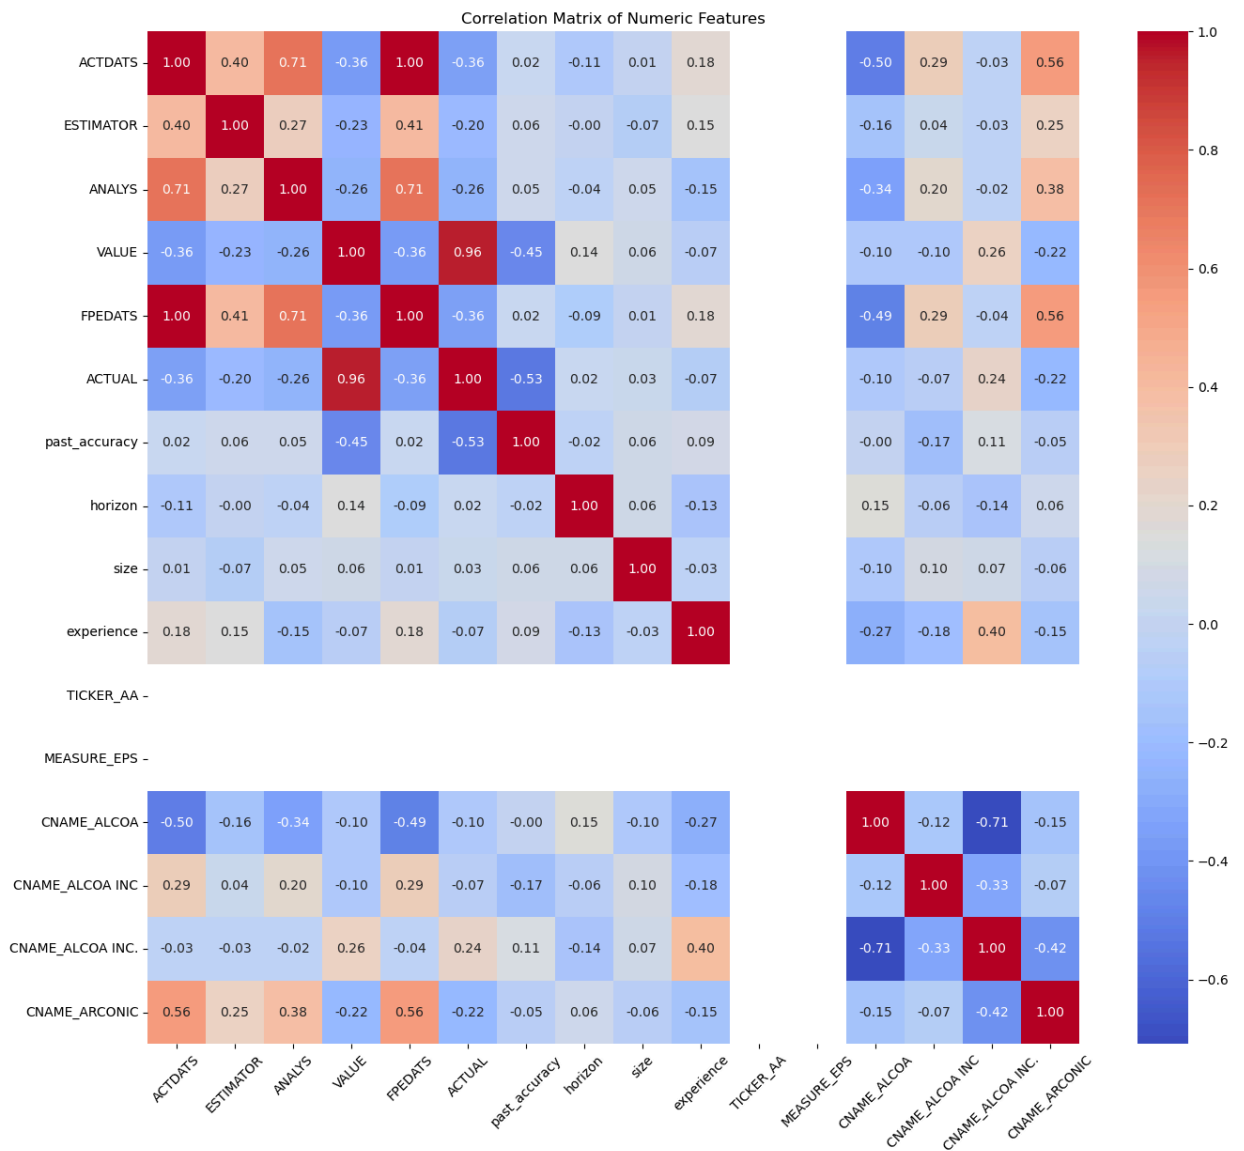


```
In [20]: # Apply one-hot encoding
clean_data = pd.get_dummies(clean_data, columns=['TICKER', 'MEASURE', 'CNAME'])
```

```
In [21]: # Exclude non-numeric columns explicitly
numeric_cols = clean_data.select_dtypes(include=[np.number]).drop(columns=['FPI']) # C

# Recalculate the correlation matrix for the DataFrame
correlation_matrix = numeric_cols.corr()

# Visualize the correlation matrix with annotations
plt.figure(figsize=(16, 14)) # Increase the figure size for better visibility
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", cbar=True, anr
plt.title('Correlation Matrix of Numeric Features')
plt.xticks(rotation=45) # Rotate x labels for better readability if necessary
plt.yticks(rotation=0)
plt.show()
```

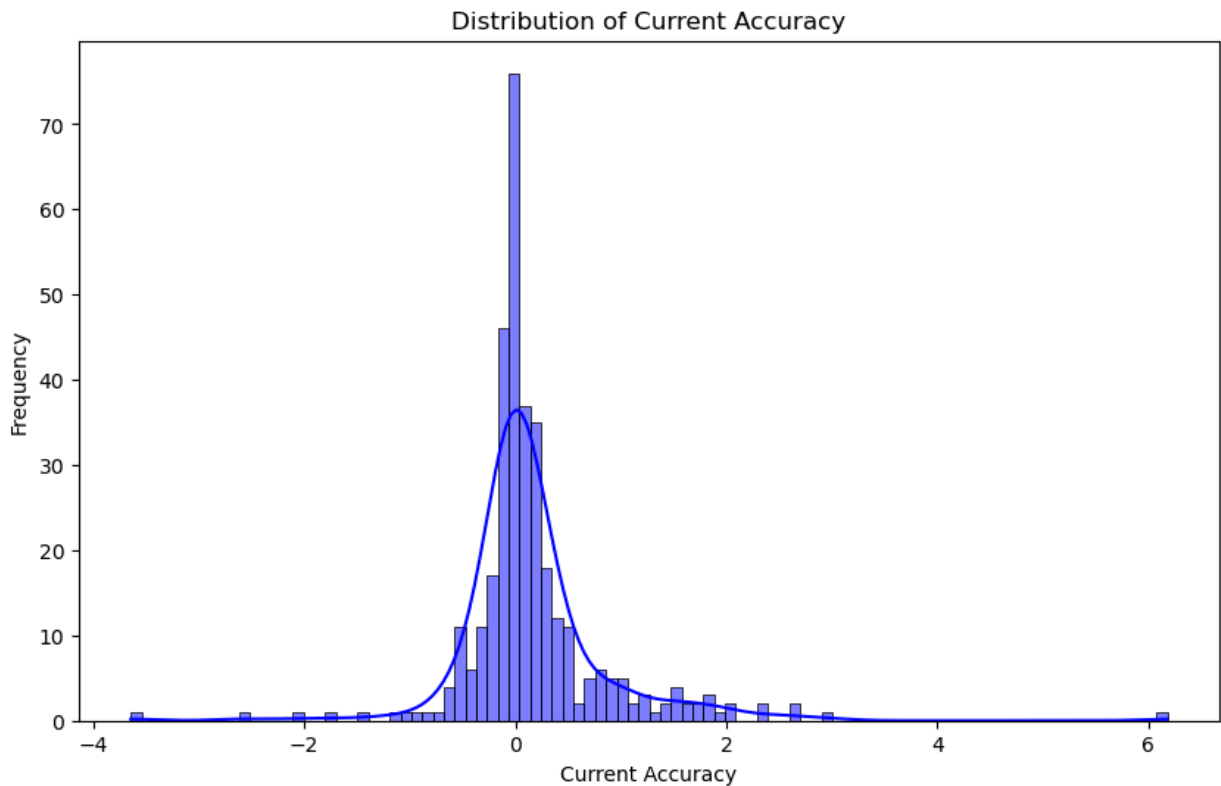


```
In [22]: # Calculate current accuracy by subtracting the actual EPS from the forecasted EPS
clean_data['current_accuracy'] = clean_data['VALUE'] - clean_data['ACTUAL']

# Display the first few rows to ensure the calculation was successful
print(clean_data[['ANALYS', 'FPEDATS', 'VALUE', 'ACTUAL', 'current_accuracy']].head())
```

	ANALYS	FPEDATS	VALUE	ACTUAL	current_accuracy
2071	127.0	1996	2.7375	2.7375	0.0000
2103	127.0	1997	3.4050	3.3375	0.0675
1544	127.0	1998	3.6000	3.6525	-0.0525
1486	127.0	2000	6.3750	5.4450	0.9300
1786	281.0	1996	4.1250	2.7375	1.3875

```
In [23]: # Plotting the distribution of current accuracy
plt.figure(figsize=(10, 6))
sns.histplot(clean_data['current_accuracy'], kde=True, color='blue')
plt.title('Distribution of Current Accuracy')
plt.xlabel('Current Accuracy')
plt.ylabel('Frequency')
plt.show()
```

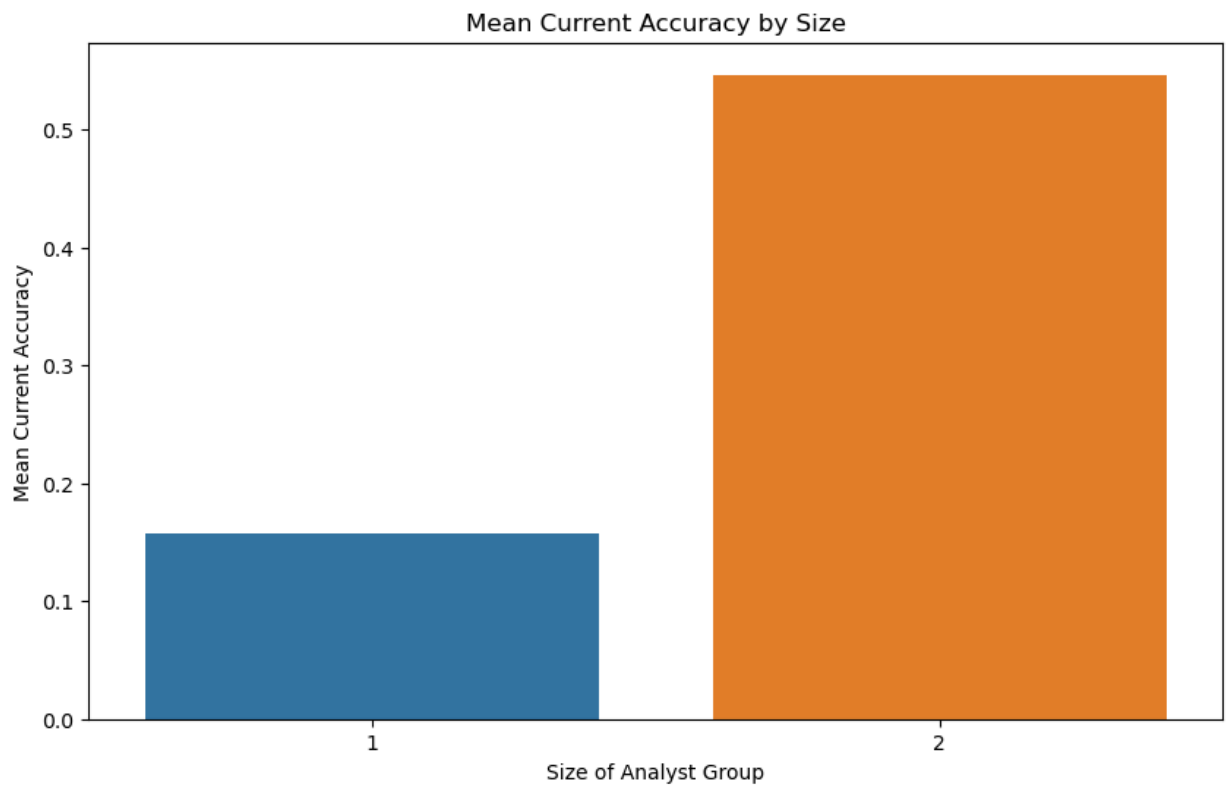
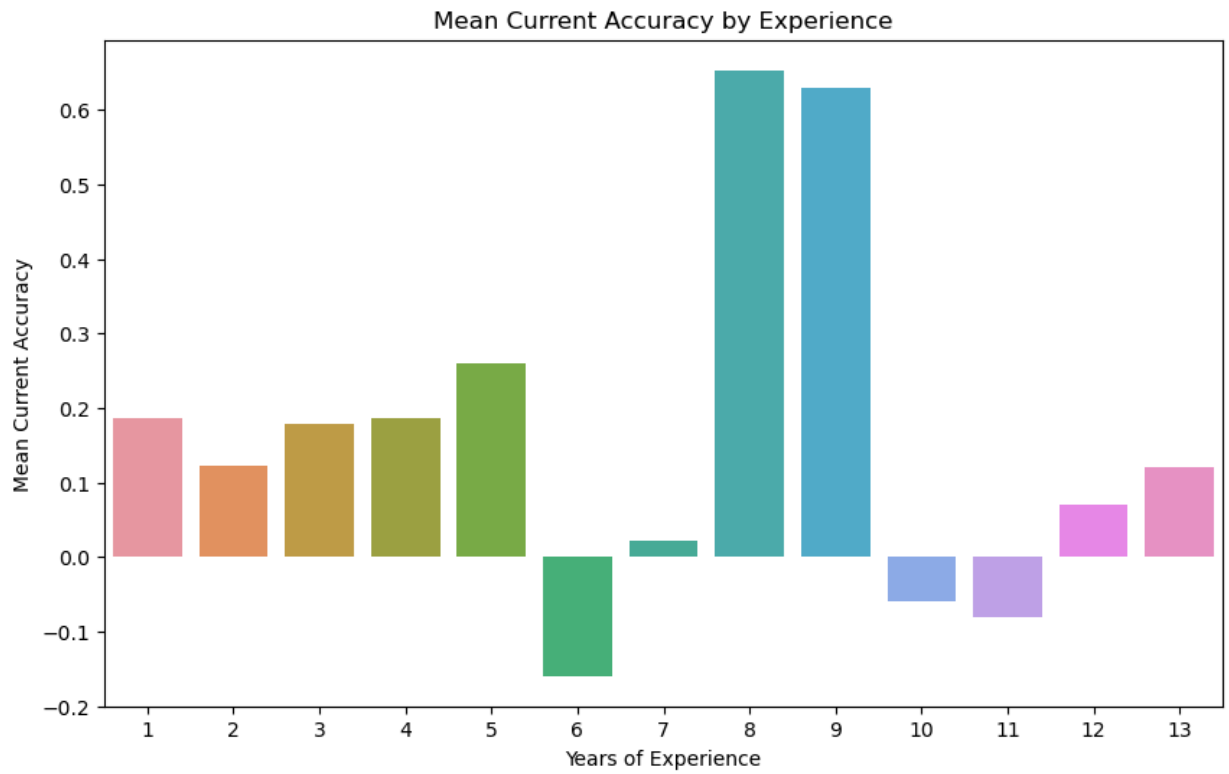


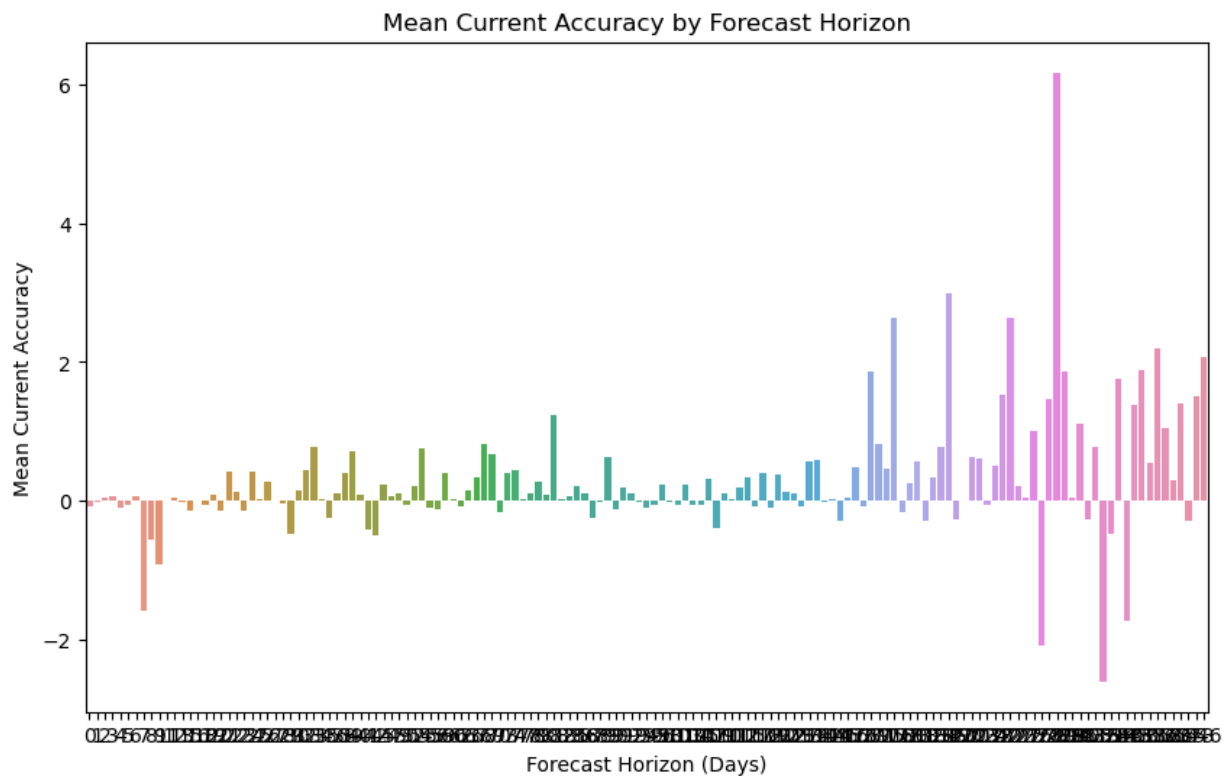
```
In [24]: # Ensure 'current_accuracy' exists and calculate mean accuracy for different groups
mean_accuracy_by_experience = clean_data.groupby('experience')['current_accuracy'].mean()
mean_accuracy_by_size = clean_data.groupby('size')['current_accuracy'].mean()
mean_accuracy_by_horizon = clean_data.groupby('horizon')['current_accuracy'].mean()

# Visualize Mean Current Accuracy by Experience
plt.figure(figsize=(10, 6))
sns.barplot(x=mean_accuracy_by_experience.index, y=mean_accuracy_by_experience.values)
plt.title('Mean Current Accuracy by Experience')
plt.xlabel('Years of Experience')
plt.ylabel('Mean Current Accuracy')
plt.show()

# Visualize Mean Current Accuracy by Size
plt.figure(figsize=(10, 6))
sns.barplot(x=mean_accuracy_by_size.index, y=mean_accuracy_by_size.values)
plt.title('Mean Current Accuracy by Size')
plt.xlabel('Size of Analyst Group')
plt.ylabel('Mean Current Accuracy')
plt.show()

# Visualize Mean Current Accuracy by Horizon
plt.figure(figsize=(10, 6))
sns.barplot(x=mean_accuracy_by_horizon.index, y=mean_accuracy_by_horizon.values)
plt.title('Mean Current Accuracy by Forecast Horizon')
plt.xlabel('Forecast Horizon (Days)')
plt.ylabel('Mean Current Accuracy')
plt.show()
```





## Splitting the dataset into TrainData and TestData

```
In [25]: # Extract the year from FPEDATS
clean_data['year'] = clean_data['FPEDATS'].astype(str).str.slice(0, 4).astype(int)
```

```
In [26]: # Find the unique years in the dataset
unique_years = clean_data['year'].unique()

# Assuming the years are sorted, find the split year
split_year = sorted(unique_years)[-4] # Get the fourth-last unique year
```

```
In [27]: # Create the train_data and test_data splits
train_data = clean_data[clean_data['year'] < split_year]
test_data = clean_data[clean_data['year'] >= split_year]
```

```
In [28]: # Verify the train_data range
train_years = train_data['year'].unique()
print(f"Training data years: {sorted(train_years)}")

# Verify the test_data range
test_years = test_data['year'].unique()
print(f"Testing data years: {sorted(test_years)}")
```

Training data years: [1996, 1997, 1998, 2000, 2001, 2002, 2003, 2004, 2005, 2006, 2007, 2008, 2009, 2010, 2011, 2012, 2014, 2015, 2016]  
 Testing data years: [2017, 2018, 2019, 2020]

```
In [29]: # Print the data types of all columns
print(train_data.dtypes)
```

```

ACTDATS                int64
ESTIMATOR              float64
ANALYS                 float64
FPI                   float64
VALUE                 float64
FPEDATS               int64
REVDATS               datetime64[ns]
REVTIMS               object
ANNDATS               datetime64[ns]
ANNTIMS               object
ACTUAL                float64
ANNDATS_ACT           datetime64[ns]
ANNTIMS_ACT           object
past_accuracy         float64
horizon               int64
size                  int64
experience             int32
TICKER_AA             uint8
MEASURE_EPS           uint8
CNAME_ALCOA           uint8
CNAME_ALCOA INC       uint8
CNAME_ALCOA INC.      uint8
CNAME_ARCONIC         uint8
current_accuracy      float64
year                  int32
dtype: object

```

```

In [30]: # Check for columns with NaN values
nan_columns = train_data.columns[train_data.isna().any()].tolist()

# Print out the columns that have NaN values
print(f"Columns with NaN values: {nan_columns}")

```

Columns with NaN values: []

# Modeling

## Backward Selection and Linear Regression

```

In [31]: import pandas as pd
import statsmodels.api as sm

# Dropping 'current_accuracy' from predictors
train_data = train_data.drop(columns=['current_accuracy'])

# Selecting only numeric columns for predictors
predictor_columns = train_data.select_dtypes(include=['int64', 'float64']).columns.drop('year')

# Prepare your predictors (X) and response (y)
X = train_data[predictor_columns]
y = train_data['ACTUAL']

# Define a function for backward elimination
def backward_elimination(X, y, significance_level=0.05):
    features = X.columns.tolist()

```

```
while len(features) > 0:
    # Fit the model with the current set of features
    X_with_constant = sm.add_constant(X[features])
    p_values = sm.OLS(y, X_with_constant).fit().pvalues[1:] # exclude the intercept
    # Find the feature with the largest p-value
    max_p_value = p_values.max()
    # If the feature has a p-value greater than the significance level, drop it
    if max_p_value >= significance_level:
        excluded_feature = p_values.idxmax()
        features.remove(excluded_feature)
    else:
        break
return features

# Perform backward elimination
selected_features = backward_elimination(X, y)

# Fit the model using the selected features
X_selected = sm.add_constant(X[selected_features])
linearModel_Team1 = sm.OLS(y, X_selected).fit()

# Print the summary of the model
print(linearModel_Team1.summary())

# Document the selected features
print("Selected features based on backward elimination:", selected_features)
```



## OLS Regression Results

Dep. Variable:	ACTUAL	R-squared:	0.940			
Model:	OLS	Adj. R-squared:	0.939			
Method:	Least Squares	F-statistic:	971.1			
Date:	Fri, 10 May 2024	Prob (F-statistic):	2.25e-187			
Time:	18:08:57	Log-Likelihood:	-292.57			
No. Observations:	317	AIC:	597.1			
Df Residuals:	311	BIC:	619.7			
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
ACTDATS	-1.714e-06	6.85e-07	-2.504	0.013	-3.06e-06	-3.67e-07
ESTIMATOR	0.0001	4.23e-05	2.422	0.016	1.92e-05	0.000
FPI	34.8524	13.742	2.536	0.012	7.813	61.891
VALUE	0.8809	0.016	53.935	0.000	0.849	0.913
past_accuracy	-0.1331	0.016	-8.175	0.000	-0.165	-0.101
horizon	-0.0027	0.000	-7.585	0.000	-0.003	-0.002
=====						
Omnibus:	78.878	Durbin-Watson:	2.013			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	2320.685			
Skew:	0.040	Prob(JB):	0.00			
Kurtosis:	16.255	Cond. No.	7.98e+09			
=====						

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 7.98e+09. This might indicate that there are strong multicollinearity or other numerical problems.

Selected features based on backward elimination: ['ACTDATS', 'ESTIMATOR', 'FPI', 'VALUE', 'past\_accuracy', 'horizon']

## KNN and Cross-Validation

```
In [32]: from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import cross_val_score
import numpy as np

import matplotlib.pyplot as plt

# Define predictors and response, excluding non-predictive and target column
X_train = train_data[predictor_columns]
y_train = train_data['ACTUAL']

# List to store average RMSE for different k values
average_rmse = []

# Try k from 1 to a reasonable upper limit, stepping by 2 (only odd k values)
k_values = range(1, 21, 2)
for k in k_values:
    knn = KNeighborsRegressor(n_neighbors=k)
    # Perform cross-validation
    scores = cross_val_score(knn, X_train, y_train, scoring='neg_mean_squared_error',
                             rmse_scores = np.sqrt(-scores))
    average_rmse.append(rmse_scores.mean())
```

```

# Determine the best k (having the lowest RMSE)
optimal_k = k_values[np.argmin(average_rmse)]
print(f"Optimal k value: {optimal_k}")

# Fit the model with optimal k
knnModel_Team1 = KNeighborsRegressor(n_neighbors=optimal_k)
knnModel_Team1.fit(X_train, y_train)
print(f"Average RMSE for each k: {list(zip(k_values, average_rmse))}")

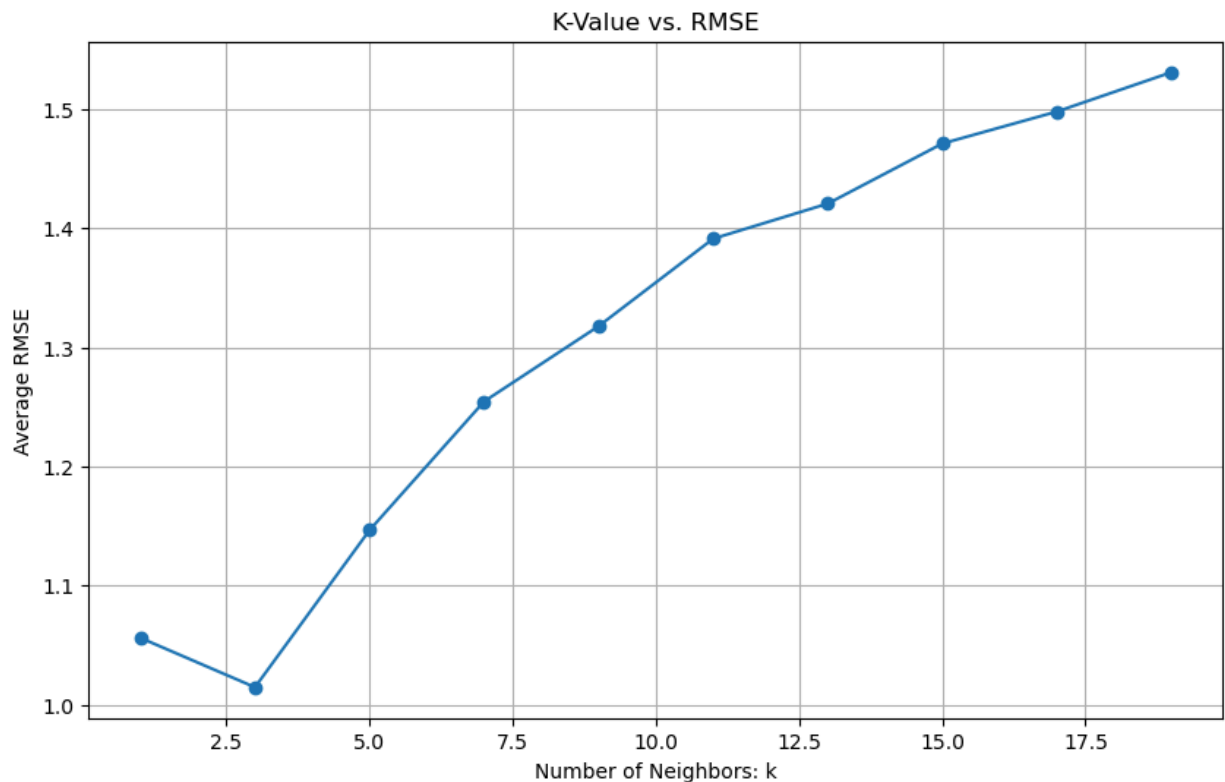
# Plot RMSE vs. k values
plt.figure(figsize=(10, 6))
plt.plot(k_values, average_rmse, marker='o')
plt.xlabel('Number of Neighbors: k')
plt.ylabel('Average RMSE')
plt.title('K-Value vs. RMSE')
plt.grid(True)
plt.show()

# Print a statement to confirm model training
print("KNN model fitted with optimal k.")

```

Optimal k value: 3

Average RMSE for each k: [(1, 1.0560580135017112), (3, 1.0145862389492148), (5, 1.1464976567517218), (7, 1.2546019907937926), (9, 1.3178889000587144), (11, 1.391343195006397), (13, 1.4207958251321862), (15, 1.4712443052994906), (17, 1.4981098226664102), (19, 1.5311610466521028)]



KNN model fitted with optimal k.

## Lasso

```

In [33]: from sklearn.linear_model import Lasso, LassoCV

# Lasso regression with cross-validation to find the best alpha
lasso = LassoCV(alphas=np.logspace(-6, 6, 13), cv=10) # Explore wide range of alpha

```

```
lasso.fit(X_train, y_train)

# Optimal alpha value
optimal_alpha = lasso.alpha_
print(f"Optimal alpha value: {optimal_alpha}")

# Final Lasso Model
lassoModel_Team1 = Lasso(alpha=optimal_alpha)
lassoModel_Team1.fit(X_train, y_train)
```

Optimal alpha value: 0.01

```
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 56.9462870418176, tolerance: 0.1927793844649123
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 56.67486979271624, tolerance: 0.1927793844649123
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 56.649002191495036, tolerance: 0.1927793844649123
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 56.647261460455645, tolerance: 0.1927793844649123
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 51.237477930617125, tolerance: 0.17949503810526318
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 51.23641460407555, tolerance: 0.17949503810526318
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 51.23778236990811, tolerance: 0.17949503810526318
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 55.542303717042806, tolerance: 0.18727504068859652
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 55.531698247235745, tolerance: 0.18727504068859652
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 55.53249697733442, tolerance: 0.18727504068859652
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 53.988634407003175, tolerance: 0.17535207887280704
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 53.81950931806262, tolerance: 0.17535207887280704
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 53.796812059935185, tolerance: 0.17535207887280704
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 53.7959640075318, tolerance: 0.17535207887280704
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 56.89640265046257, tolerance: 0.17561427623684214
    model = cd_fast.enet_coordinate_descent_gram(
```

```
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 56.6245359135943, tolerance: 0.17561427623684214
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 56.59865795039713, tolerance: 0.17561427623684214
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 56.597271152481106, tolerance: 0.17561427623684214
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 56.05041860926797, tolerance: 0.1783419930175439
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 55.769082351423165, tolerance: 0.1783419930175439
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 55.74210780779696, tolerance: 0.1783419930175439
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 55.74060260653126, tolerance: 0.1783419930175439
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 55.04041654640675, tolerance: 0.16708272762280704
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 54.75291064862857, tolerance: 0.16708272762280704
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 54.725142612838766, tolerance: 0.16708272762280704
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 54.72317922908496, tolerance: 0.16708272762280704
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 39.45239015764131, tolerance: 0.16889968117132867
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 39.18842705088936, tolerance: 0.16889968117132867
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 39.16109001069281, tolerance: 0.16889968117132867
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the n
umber of iterations. Duality gap: 39.1592681992222, tolerance: 0.16889968117132867
    model = cd_fast.enet_coordinate_descent_gram(
```

```

C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 50.90760547877037, tolerance: 0.1582522692832168
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 48.48228066561619, tolerance: 0.1582522692832168
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 48.17783086576763, tolerance: 0.1582522692832168
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 48.148241114301996, tolerance: 0.1582522692832168
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 48.14600568711771, tolerance: 0.1582522692832168
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 54.12293079154994, tolerance: 0.17314835795454547
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 53.84539755184933, tolerance: 0.17314835795454547
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 53.81879202023219, tolerance: 0.17314835795454547
    model = cd_fast.enet_coordinate_descent_gram(
C:\Users\YOG\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
617: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 53.816976316168436, tolerance: 0.17314835795454547
    model = cd_fast.enet_coordinate_descent_gram(

```

Out[33]:

▼ Lasso

Lasso(alpha=0.01)

## Random Forest

In [34]:

```

from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV

# Set up the parameter grid
param_grid = {'max_features': range(1, len(predictor_columns)+1)}

# Create Random Forest model
rf = RandomForestRegressor(n_estimators=100, random_state=42)

# Conduct grid search
grid_search = GridSearchCV(rf, param_grid, cv=5, scoring='neg_mean_squared_error')
grid_search.fit(X_train, y_train)

# Best max_features
optimal_max_features = grid_search.best_params_['max_features']

```

```
print(f"Optimal max_features: {optimal_max_features}")

# Final Random Forest Model
RFModel_Team1 = RandomForestRegressor(n_estimators=100, max_features=optimal_max_featu
RFModel_Team1.fit(X_train, y_train)
```

Optimal max\_features: 3

Out[34]:

```
RandomForestRegressor
RandomForestRegressor(max_features=3, random_state=42)
```

## Calculating MSPE, MSE, r2\_SCROCE, MAE Matrix

In [35]:

```
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

# Prepare test data
X_test = test_data[predictor_columns]
y_test = test_data['ACTUAL']

# Predict using each model
y_pred_linear = linearModel_Team1.predict(sm.add_constant(X_test[selected_features]))
y_pred_knn = knnModel_Team1.predict(X_test)
y_pred_lasso = lassoModel_Team1.predict(X_test)
y_pred_rf = RFModel_Team1.predict(X_test)

# Calculate RMSE for each model
rmse_linear = np.sqrt(mean_squared_error(y_test, y_pred_linear))
rmse_knn = np.sqrt(mean_squared_error(y_test, y_pred_knn))
rmse_lasso = np.sqrt(mean_squared_error(y_test, y_pred_lasso))
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))

# Calculate MAE for each model
mae_linear = mean_absolute_error(y_test, y_pred_linear)
mae_knn = mean_absolute_error(y_test, y_pred_knn)
mae_lasso = mean_absolute_error(y_test, y_pred_lasso)
mae_rf = mean_absolute_error(y_test, y_pred_rf)

# Calculate R-squared for each model
r2_linear = r2_score(y_test, y_pred_linear)
r2_knn = r2_score(y_test, y_pred_knn)
r2_lasso = r2_score(y_test, y_pred_lasso)
r2_rf = r2_score(y_test, y_pred_rf)

# Calculate MSE for each model
mse_linear = mean_squared_error(y_test, y_pred_linear)
mse_knn = mean_squared_error(y_test, y_pred_knn)
mse_lasso = mean_squared_error(y_test, y_pred_lasso)
mse_rf = mean_squared_error(y_test, y_pred_rf)

# Calculate Squared Percentage Error for each model
def calculate_mspe(y_true, y_pred):
```



```

# Handling division by zero or very small values in actual by adding a small number
epsilon = np.finfo(float).eps
squared_percentage_errors = ((y_true - y_pred) / (y_true + epsilon)) ** 2
return np.mean(squared_percentage_errors)

# Calculating MSPE for each model
mspe_linear = calculate_mspe(y_test, y_pred_linear)
mspe_knn = calculate_mspe(y_test, y_pred_knn)
mspe_lasso = calculate_mspe(y_test, y_pred_lasso)
mspe_rf = calculate_mspe(y_test, y_pred_rf)

# Print the metrics
print("RMSE (linear, KNN, Lasso, RF):", rmse_linear, rmse_knn, rmse_lasso, rmse_rf)
print("MAE (linear, KNN, Lasso, RF):", mae_linear, mae_knn, mae_lasso, mae_rf)
print("R-squared (linear, KNN, Lasso, RF):", r2_linear, r2_knn, r2_lasso, r2_rf)
print("MSE (Linear, KNN, Lasso, RF):", mse_linear, mse_knn, mse_lasso, mse_rf)
print("MSPE (Linear, KNN, Lasso, RF):", mspe_linear, mspe_knn, mspe_lasso, mspe_rf)

RMSE (linear, KNN, Lasso, RF): 0.22604471848830784 0.6712908082242368 0.2367711940393
287 0.28769123973392535
MAE (linear, KNN, Lasso, RF): 0.18075039834991805 0.5660714285714286 0.19971369743977
305 0.17233214285714274
R-squared (linear, KNN, Lasso, RF): 0.7613798446560578 -1.1044557422304555 0.73819609
25055256 0.6134802668282902
MSE (Linear, KNN, Lasso, RF): 0.05109621475645835 0.45063134920634906 0.0560605983268
0944 0.0827662494196429
MSPE (Linear, KNN, Lasso, RF): 0.03564689507104211 0.1555992018785088 0.0368576359558
7114 0.11638758849118973

```

```

In [36]: # Create a DataFrame to store the results
model_results = pd.DataFrame({
    'Model': ['Linear Regression', 'KNN', 'Lasso', 'Random Forest'],
    'RMSE': [rmse_linear, rmse_knn, rmse_lasso, rmse_rf],
    'MAE': [mae_linear, mae_knn, mae_lasso, mae_rf],
    'R-squared': [r2_linear, r2_knn, r2_lasso, r2_rf],
    'MSE': [mse_linear, mse_knn, mse_lasso, mse_rf],
    'MSPE': [mspe_linear, mspe_knn, mspe_lasso, mspe_rf]
})

# Print the results DataFrame
print(model_results.to_string())

```

	Model	RMSE	MAE	R-squared	MSE	MSPE
0	Linear Regression	0.226045	0.180750	0.761380	0.051096	0.035647
1	KNN	0.671291	0.566071	-1.104456	0.450631	0.155599
2	Lasso	0.236771	0.199714	0.738196	0.056061	0.036858
3	Random Forest	0.287691	0.172332	0.613480	0.082766	0.116388

```

In [37]: # Convert the 'Model' column to categorical data type
model_results['Model'] = model_results['Model'].astype('category')

# Find the best model based on the 'MSPE' column
best_model = model_results.loc[model_results['MSPE'].idxmin()]

# Print the best model
print(f"Best Model: {best_model['Model']}")

```

Best Model: Linear Regression



# MSPE of Consensus

```
In [38]: # Assuming 'test_data' contains columns 'year' and 'ACTUAL' where 'ACTUAL' represents
consensus_forecast = test_data.groupby('year')['ACTUAL'].mean().reset_index(name='ConsensusForecast')
print(consensus_forecast)

# Merge the consensus forecast back to the test data using the 'year' column as a key
test_data_with_consensus = test_data.merge(consensus_forecast, on='year', how='left')
print(test_data_with_consensus.head())
```

	year	ConsensusForecast
0	2017	1.22
1	2018	1.36
2	2019	2.11
3	2020	0.80

	ACTDATS	ESTIMATOR	ANALYS	FPI	VALUE	FPEDATS	REVDATS	REVTIMS	\
0	20181102	11.0	31736.0	1.0	1.33	2018	2019-01-11	12:23:56	
1	20191126	11.0	31736.0	1.0	2.11	2019	2020-01-10	8:03:53	
2	20181030	282.0	43401.0	1.0	1.30	2018	2018-11-13	6:45:10	
3	20171023	157.0	73867.0	1.0	1.13	2017	2017-12-01	18:05:27	
4	20181220	2301.0	77011.0	1.0	1.30	2018	2018-12-20	0:03:57	

	ANNDATS	ANNTIMS	...	experience	TICKER_AA	MEASURE_EPS	CNAME_ALCOA	\
0	2018-11-02	8:04:00	...		1	1	0	
1	2019-11-26	8:01:00	...		2	1	0	
2	2018-10-30	16:17:00	...		1	1	0	
3	2017-10-23	22:25:00	...		1	1	0	
4	2018-12-20	0:02:00	...		1	1	0	

	CNAME_ALCOA	INC	CNAME_ALCOA	INC.	CNAME_ARCONIC	current_accuracy	year	\
0		0		0	1	-0.03	2018	
1		0		0	1	0.00	2019	
2		0		0	1	-0.06	2018	
3		0		0	1	-0.09	2017	
4		0		0	1	-0.06	2018	

	ConsensusForecast
0	1.36
1	2.11
2	1.36
3	1.22
4	1.36

[5 rows x 26 columns]

```
In [39]: # Calculate MSPE for the consensus forecast
mspe_consensus = ((test_data_with_consensus['ACTUAL'] - test_data_with_consensus['ConsensusForecast'])**2).mean()
print(f"Consensus Forecast MSPE: {mspe_consensus}")
```

Consensus Forecast MSPE: 2.0634963689863236e-33