

Yog Chaudhary

11727095

ADTA 5240 Week 3'rd (harvesting, Storing, And Retrieving Data)

Professor: Dr. Zeynep Orhan

Sep 15, 2023

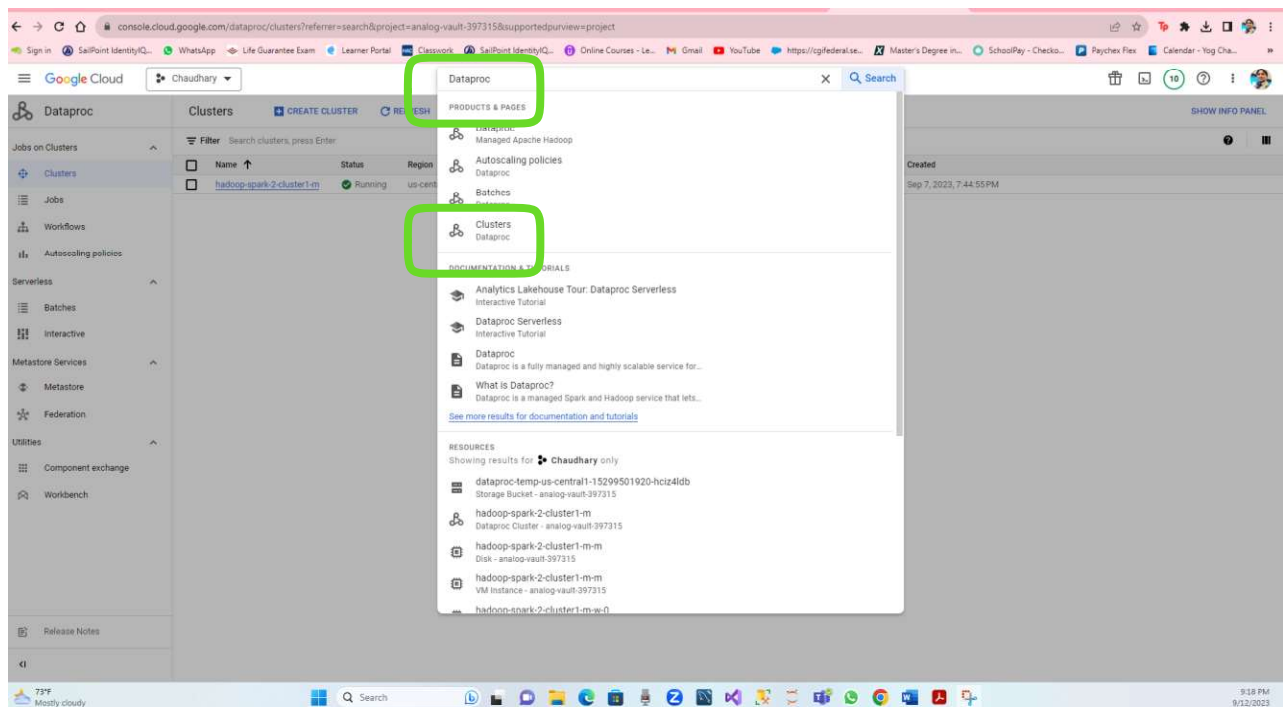
University Of North Texas

Exploring Remote Virtual Machine in the Cloud. Hadoop Ecosystem with Simple Linux Commands. Pdf

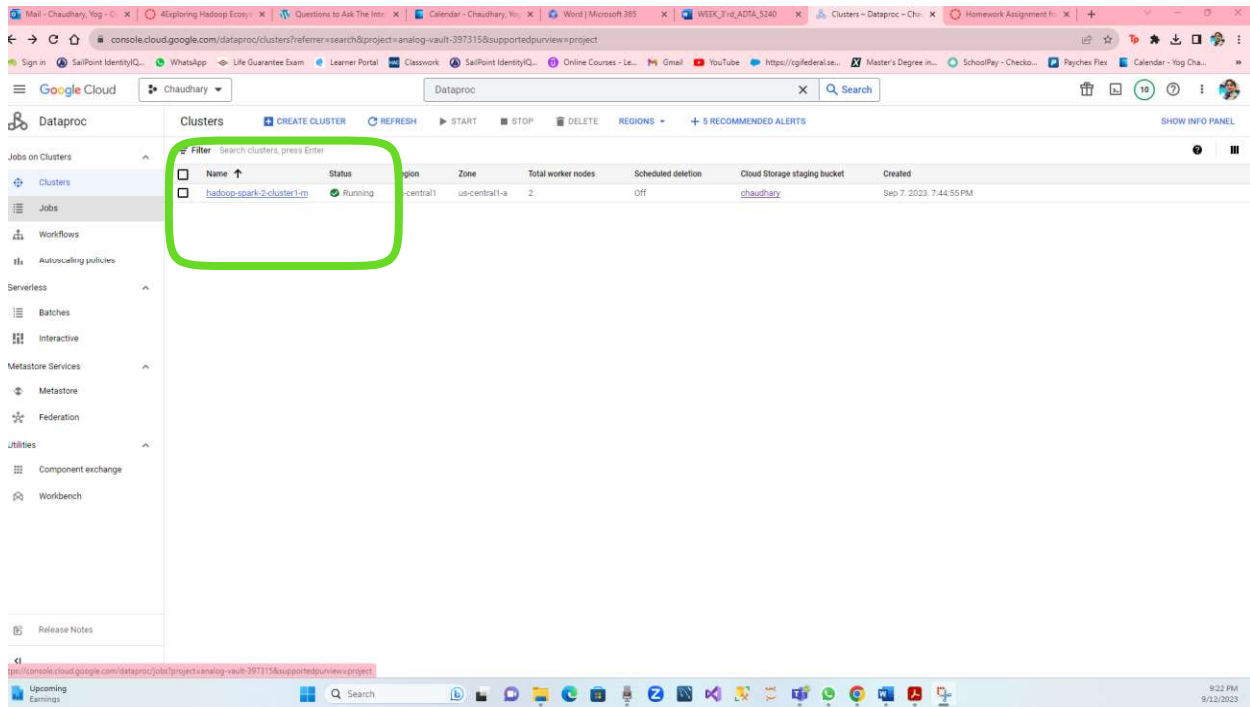
❖ Monitoring the cluster in GCP:

1. For a Google console.

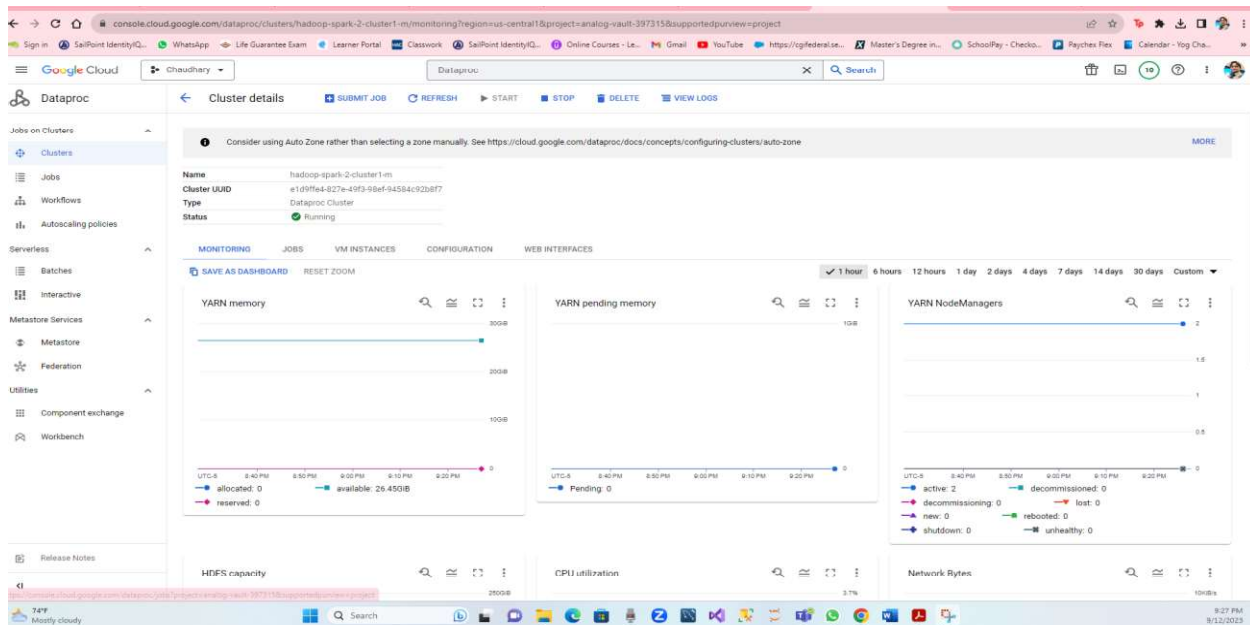
- Then I clicked on three horizontal lines.
- In the search bar I typed “Cluster” and clicked on “Clusters Dataproc.”
- Here is screenshot



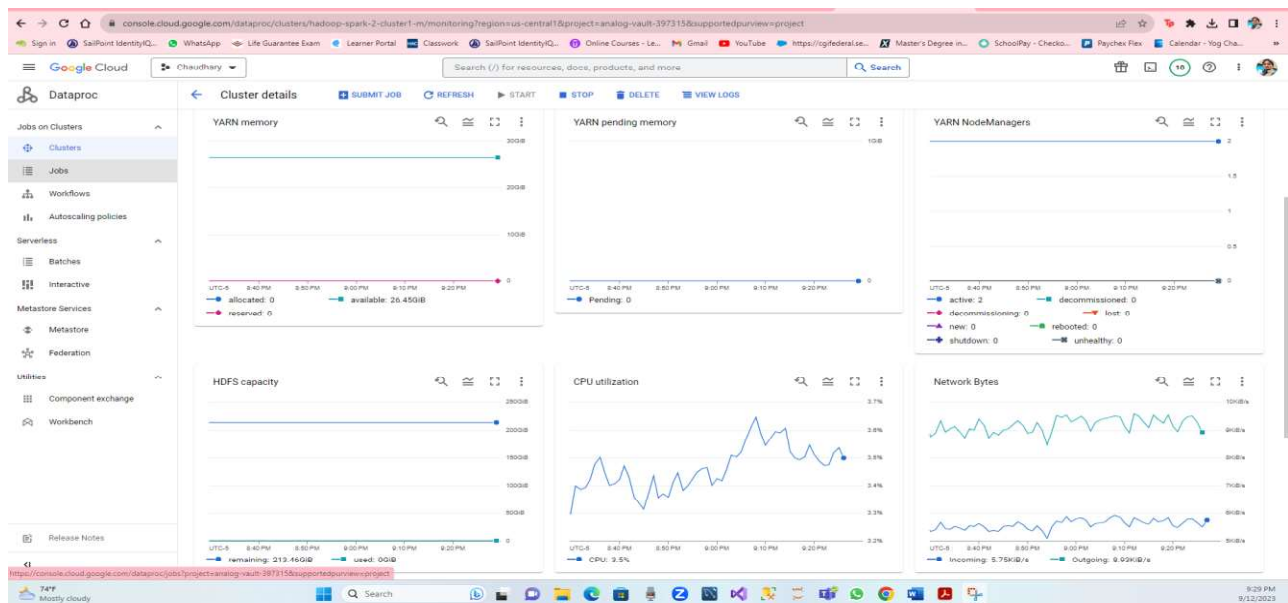
- After clicking on clusters data proc it will mention the cluster that I have previously created, and it will show that it was running (with a green check mark).
- Below shows a screenshot of that.



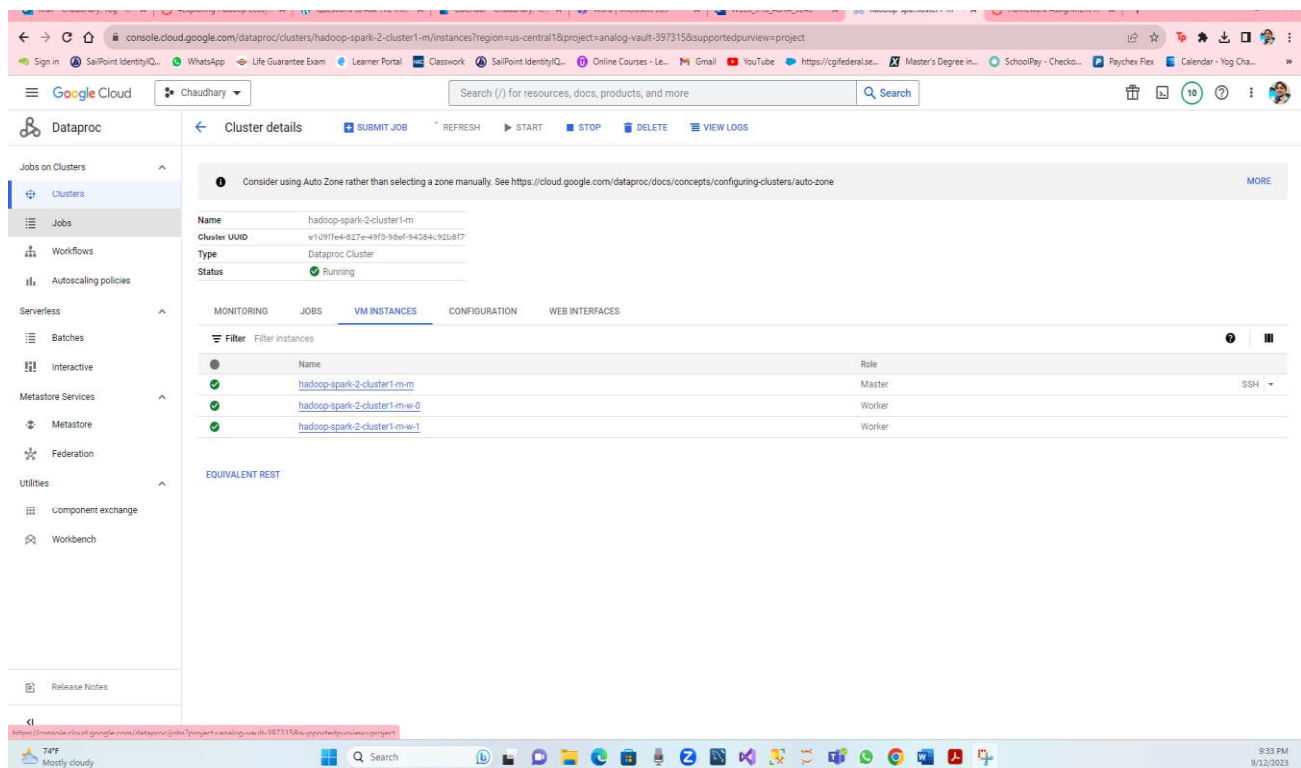
- While doing this we need to start nodes by checking on the navigation panel.
- Clicked on computing engine.
- Then clicked on three vertical dots and clicked on start/resume.
- Now click.
- Below screenshot of that.



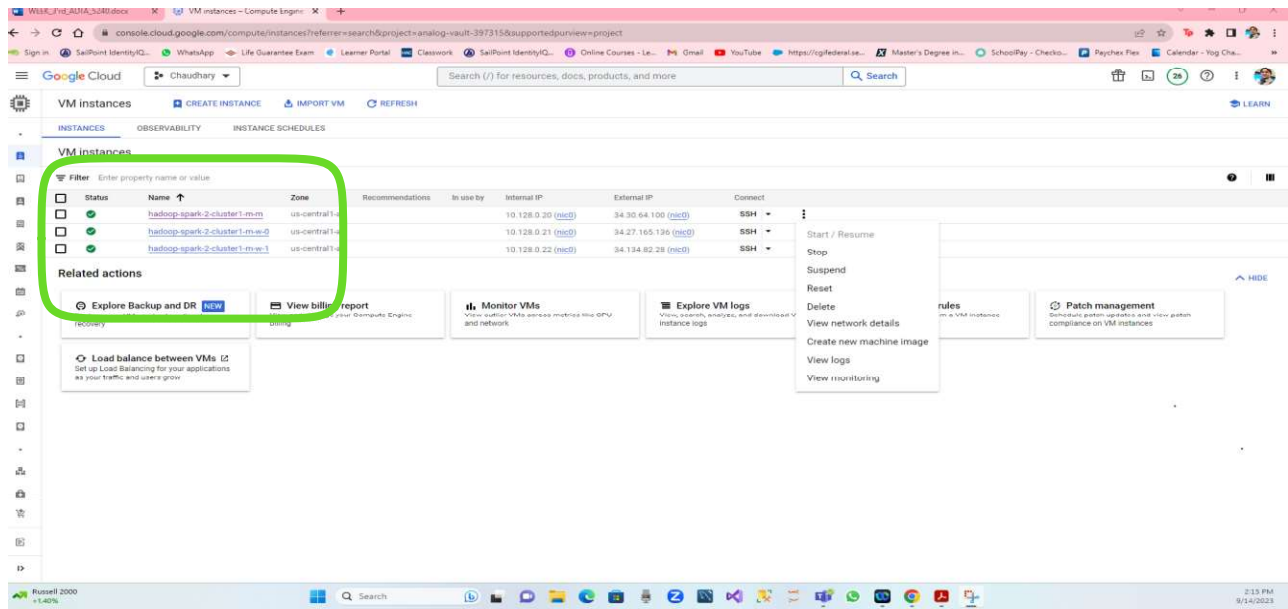
- When I scroll down it will show how the monitoring is changing according to the usage.



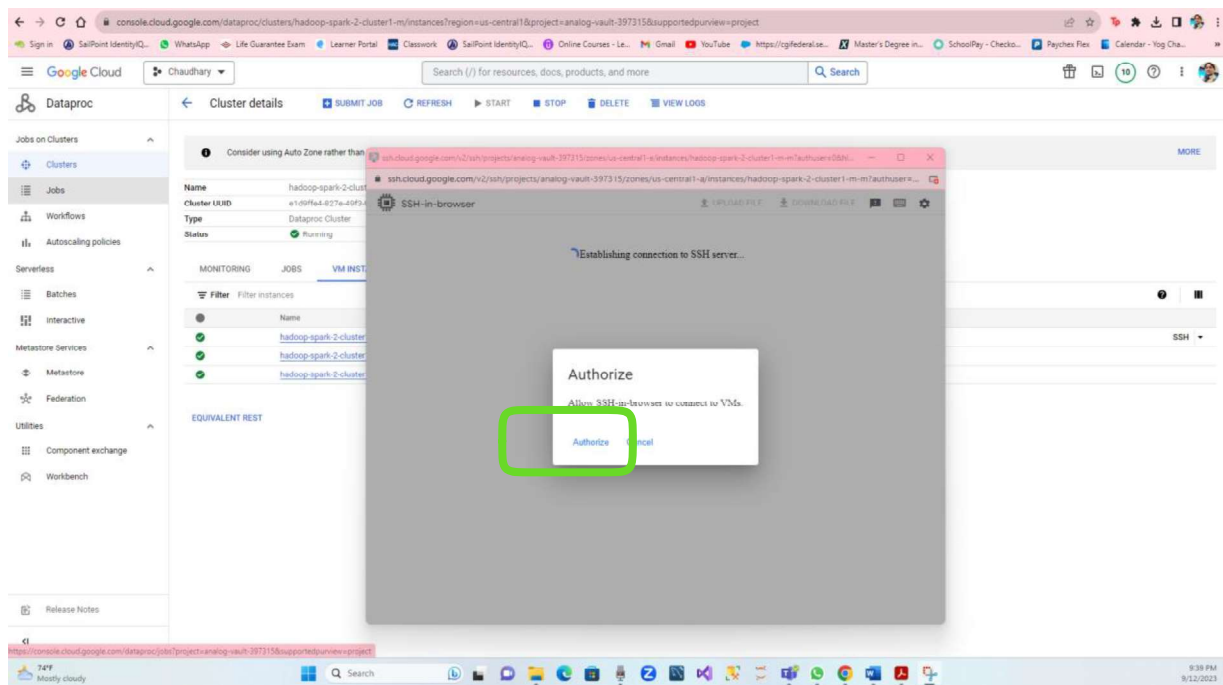
- By scrolling up and clicking on the virtual machines the screen will show the cluster details.
- We will see one master node and one worker node.
- After that I accessed the master node through “SSH.”
- Click on “SSH.”
- Click on “Open in browser.”



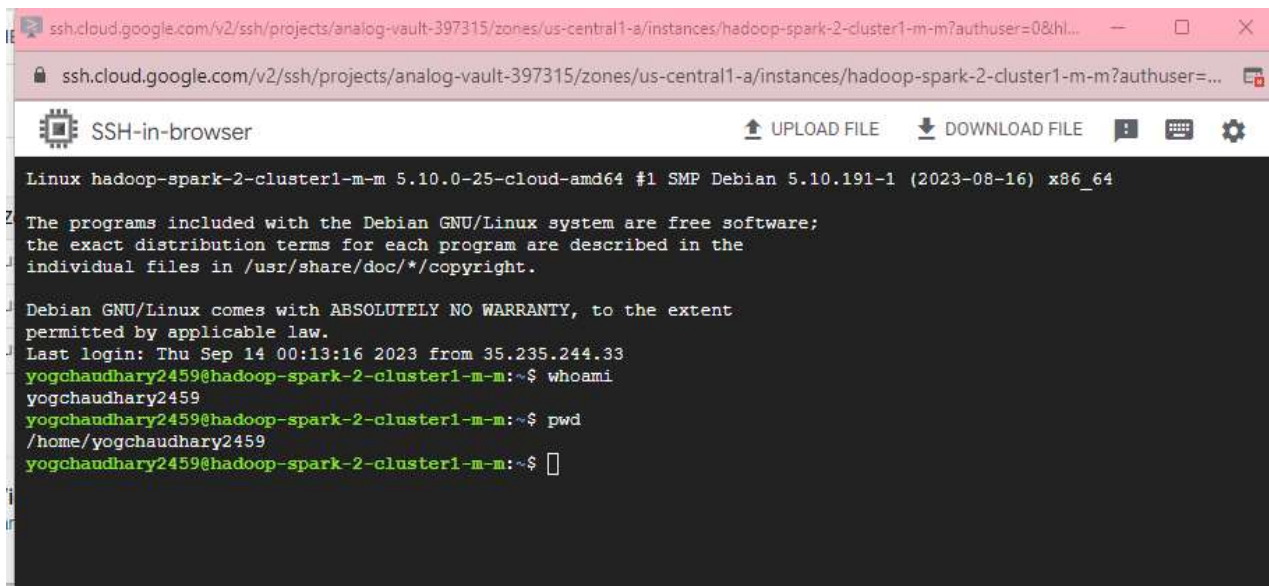
- Click on the drop-down button next to “SSH”.
- Click on open in a new browser window.



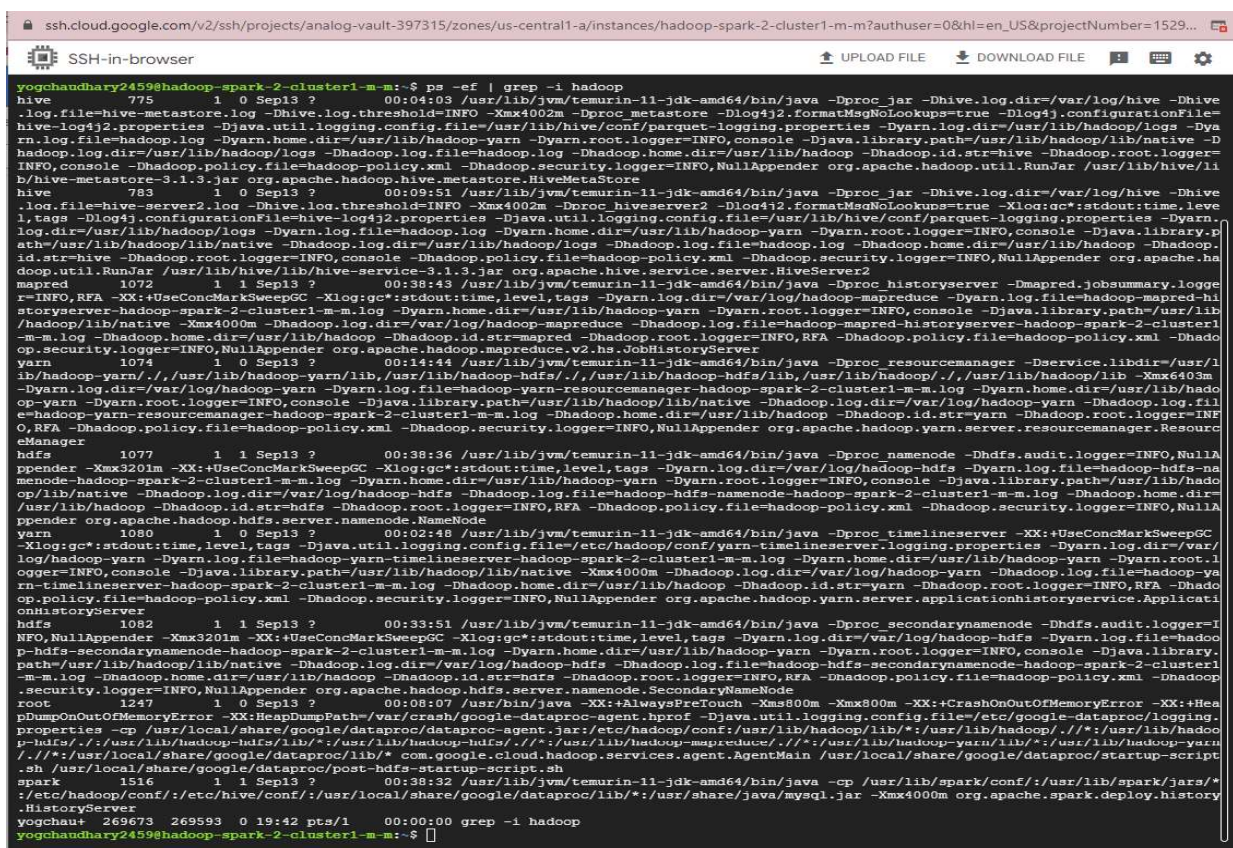
- Click authorized.



- ❖ See all the services of Hadoop in our cluster.
- ❖ Use the command.
 - whoami
 - pwd



- ❖ use other commands.
- `ps -ef | grep -I hadoop`
- Hare screenshot



- Click on “Show Scrollbar” to see the scrollbar.

Hadoop with Linux Commands is successfully screenshot

2. Now we must open another new SSH terminal.

- Going back to GCP and typing “computer engine.”
- New clicked on “SSH” then I clicked on “open in browser.”

The screenshot displays the Google Cloud Platform console interface. At the top, the navigation bar includes the Google Cloud logo, the user's name 'Chaudhary', and a search bar. Below this, the 'VM instances' section is active, showing a list of instances. The table below lists three instances, all with a status of 'Running' and located in the 'us-central1-a' zone. Each instance has an internal IP, an external IP, and an SSH connection option.

| Status | Name | Zone | Internal IP | External IP | Connect |
|---------|-------------------------------|---------------|-------------|---------------|---------|
| Running | hadoop-spark-2-cluster1-m-m | us-central1-a | 10.128.0.20 | 34.30.64.100 | SSH |
| Running | hadoop-spark-2-cluster1-m-w-0 | us-central1-a | 10.128.0.21 | 34.27.165.136 | SSH |
| Running | hadoop-spark-2-cluster1-m-w-1 | us-central1-a | 10.128.0.22 | 34.134.82.28 | SSH |

Below the table, there are several 'Related actions' cards, including 'Explore Backup and DR', 'View billing report', 'Monitor VMs', 'Explore VM logs', 'Set up firewall rules', and 'Patch management'. The bottom of the screen shows a Windows taskbar with various application icons and the system clock indicating 3:47 PM on 9/14/2023.


```
Linux hadoop-spark-2-cluster1-m-w-0 5.10.0-25-cloud-amd64 #1 SMP Debian 5.10.191-1 (2023-08-16) x86_64

The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.
Last login: Wed Sep 13 05:29:51 2023 from 35.235.244.33
yogchaudhary2459@hadoop-spark-2-cluster1-m-w-0:~$ whoami
yogchaudhary2459
yogchaudhary2459@hadoop-spark-2-cluster1-m-w-0:~$ pwd
/home/yogchaudhary2459
yogchaudhary2459@hadoop-spark-2-cluster1-m-w-0:~$ ps -ef | grep -i hadoop
hdfs      1097      1  0 Sep13 ?        00:05:13 /usr/lib/jvm/temurin-11-jdk-amd64/bin/java -Dproc_datanode
-Xmx512m -Dyarn.log.dir=/var/log/hadoop-hdfs -Dyarn.log.file=hadoop-hdfs-datanode-hadoop-spark-2-cluster1-m-w-0
.log -Dyarn.home.dir=/usr/lib/hadoop-yarn -Dyarn.root.logger=INFO,console -Djava.library.path=/usr/lib/hadoop/lib/
native -Dhadoop.log.dir=/var/log/hadoop-hdfs -Dhadoop.log.file=hadoop-hdfs-datanode-hadoop-spark-2-cluster1-
m-w-0.log -Dhadoop.home.dir=/usr/lib/hadoop -Dhadoop.id.str=hdfs -Dhadoop.root.logger=INFO,RFA -Dhadoop.policy.
file=hadoop-policy.xml -Dhadoop.security.logger=INFO,NullAppender org.apache.hadoop.hdfs.server.datanode.DataNo
de
yarn      1108      1  1 Sep13 ?        00:40:00 /usr/lib/jvm/temurin-11-jdk-amd64/bin/java -Dproc_nodemanager
-Dyarn.log.dir=/var/log/hadoop-yarn -Dyarn.log.file=hadoop-yarn-nodemanager-hadoop-spark-2-cluster1-m-w-0.lo
g -Dyarn.home.dir=/usr/lib/hadoop-yarn -Dyarn.root.logger=INFO,console -Djava.library.path=/usr/lib/hadoop/lib/
native -Xmx1638m -Dhadoop.log.dir=/var/log/hadoop-yarn -Dhadoop.log.file=hadoop-yarn-nodemanager-hadoop-spark-2
-cluster1-m-w-0.log -Dhadoop.home.dir=/usr/lib/hadoop -Dhadoop.id.str=yarn -Dhadoop.root.logger=INFO,RFA -Dhado
op.policy.file=hadoop-policy.xml -Dhadoop.security.logger=INFO,NullAppender org.apache.hadoop.yarn.server.nodem
anager.NodeManager
root      1191      1  0 Sep13 ?        00:03:37 /usr/bin/java -XX:+AlwaysPreTouch -Xms150m -Xmx150m -XX:+Cr
ashOnOutOfMemoryError -XX:+HeapDumpOnOutOfMemoryError -XX:HeapDumpPath=/var/crash/google-dataproc-agent.hprof -
Djava.util.logging.config.file=/etc/google-dataproc/logging.properties -cp /usr/local/share/google/dataproc/dat
aproc-agent.jar:/etc/hadoop/conf:/usr/lib/hadoop/lib/*:/usr/lib/hadoop/./*/:/usr/lib/hadoop-hdfs/./*/:/usr/lib/ha
dooop-hdfs/lib/*:/usr/lib/hadoop-hdfs/./*/:/usr/lib/hadoop-mapreduce/./*/:/usr/lib/hadoop-yarn/lib/*:/usr/lib/ha
dooop-yarn/./*/:/usr/local/share/google/dataproc/lib/* com.google.cloud.hadoop.services.agent.AgentMain /usr/loc
al/share/google/dataproc/startup-script.sh /usr/local/share/google/dataproc/post-hdfs-startup-script.sh
yogchau+  39496  39478  0 20:43 pts/0    00:00:00 grep -i hadoop
yogchaudhary2459@hadoop-spark-2-cluster1-m-w-0:~$ ls -
ls: cannot access '-': No such file or directory
yogchaudhary2459@hadoop-spark-2-cluster1-m-w-0:~$
```

- Now finally all 3 virtual machine instances were stopped as shown below in GCP by selecting stop from the three dots present at the top right to the **SSH** each node.

console.cloud.google.com/compute/instances?project=analog-vault-397315&supportedpurview=project

Google Cloud Chaudhary Search (/) for resources, docs, products, and more

VM instances CREATE INSTANCE IMPORT VM REFRESH

INSTANCES OBSERVABILITY INSTANCE SCHEDULES

VM instances

Filter Enter property name or value

| Status | Name | Zone | Recommendations | In use by | Internal IP | External IP | Connect |
|--------------------------|-------------------------------|---------------|-----------------|-----------|--------------------|-------------|---------|
| <input type="checkbox"/> | hadoop-spark-2-cluster1-m-m | us-central1-a | | | 10.128.0.20 (nic0) | | SSH |
| <input type="checkbox"/> | hadoop-spark-2-cluster1-m-w-0 | us-central1-a | | | 10.128.0.21 (nic0) | | SSH |
| <input type="checkbox"/> | hadoop-spark-2-cluster1-m-w-1 | us-central1-a | | | 10.128.0.22 (nic0) | | SSH |

Related actions

- Explore Backup and DR NEW: Back up your VMs and set up disaster recovery
- View billing report: View and manage your Compute Engine billing
- Monitor VMs: View outlier VMs across metrics like CPU and network
- Explore VM logs: View, search, analyze, and download VM instance logs
- Set up firewall rules: Control traffic to and from a VM instance
- Patch management: Schedule patch updates and view patch compliance on VM instances

4. The components of the Ecosystem for record and steps, included screenshot
 → Here is the screenshot

